

Core Sheet Model — Architecture and Plan

Context

- Runs client-first inside a Next.js app route (client component).
- Real-time collaboration via Yjs over a Node WebSocket server.
- Persistence to PostgreSQL (event log + periodic snapshots).
- Rendering with Canvas2D initially; WebGL can be introduced later without changing the model.
- Minimal server logic; compute is local unless explicitly offloaded.

Primary Responsibilities

- Hold a declarative, immutable-ish model of a “Sheet” (board) as the single source of truth.
 - Provide deterministic command processing (undo/redo).
 - Maintain derived indices for speed (spatial index, z-order, anchors).
 - Enable artifact extensibility (MDX, Mermaid, TaskList, Box, Connector, etc.).
 - Support presence and multi-user merges via CRDT.
 - Stream exports (SVG/PNG/JSON) without coupling to UI.
-

Functional Scope (Stage)

- Create/move/resize/delete artifacts.
 - Connect artifacts with orthogonal connectors.
 - Snap to grid, edges, centers.
 - Grouping/layers/z-order.
 - Multi-select and marquee.
 - MDX and Mermaid artifacts (rendered client-side).
 - Export to PNG/SVG.
 - Real-time cursors and selections.
 - Postgres persistence with snapshots.
-

Minimal Data Structures (Essential Types)

```
type Id = string;

type Vec2 = { x: number; y: number };
type AABB = { x: number; y: number; w: number; h: number }; // axis-aligned

type Sheet = {
  id: Id;
  name: string;
  config: { grid: number; snap: boolean; theme: 'light'|'dark' };
  // Normalized stores
  artifacts: Record<Id, Artifact>;
```

```

edges: Record<Id, Edge>;
order: Id[]; // z-order, back-front (only artifact ids)
// Ephemeral derived indexes (not persisted, recomputed)
_derived?: {
  spatial: QuadtreeIndex;
  anchors: AnchorIndex;
  selection: Set<Id>;
  dirtyAABBs: AABB[];
};
};

type Artifact =
| { id: Id; type: 'box'; geom: AABB; style: Style; props: { label?: string } }
| { id: Id; type: 'mdx'; geom: AABB; style: Style; props: { source: string } }
| { id: Id; type: 'mermaid'; geom: AABB; style: Style; props: { source: string } }
| { id: Id; type: 'tasklist'; geom: AABB; style: Style; props: { items: TaskItem[] } };

type Edge = {
  id: Id;
  from: { id: Id; port?: string };
  to: { id: Id; port?: string };
  type: 'ortho';
  props?: { labels?: string[] };
};

type Style = { fill?: string; stroke?: string; strokeWidth?: number; radius?: number };

type Event = { ts: number; actor: Id; type: string; payload: unknown }; // appended to event log

// Command -> Reducer op
type Command =
| { type: 'ART_CREATE'; art: Artifact }
| { type: 'ART_PATCH'; id: Id; geom?: AABB; style?: Style; props?: any }
| { type: 'ART_DELETE'; id: Id }
| { type: 'EDGE_CREATE'; edge: Edge }
| { type: 'EDGE_DELETE'; id: Id }
| { type: 'ORDER_BRING_FRONT'; id: Id }
| { type: 'ORDER_SEND_BACK'; id: Id }
| { type: 'SELECT_SET'; ids: Id[] }
| { type: 'SELECT_ADD'; ids: Id[] }
| { type: 'SELECT_CLEAR' }
| { type: 'SHEET_CONFIG'; config: Partial<Sheet['config']> };

```

Notes

- Artifact is a tagged union. New types extend by adding a case and a schema.
- Sheet._derived is ephemeral; rebuilt after each reducer tick.
- Geometry is world-space AABB. Rotation deferred out of stage scope to keep router and snap simple.

CRDT + Event Sourcing

- CRDT document per sheet: $Y.Map()$ containing maps for artifacts, edges, order, config.
 - Local commands reduce into CRDT changes (deterministic mapping).
 - The server stores incoming CRDT updates into an append-only events table and periodically materializes snapshots.
 - Undo/redo on client as a command stack; CRDT stays canonical across peers.
-

Derived Indices

Spatial Index (Quadtree)

Purpose: $O(\log n)$ insert/search for hit-testing, marquee selection, and snap candidate lookup.

- Node stores a bounded array of entries $\{ id, aabb \}$.
- Split threshold K (e.g., 8).
- Rebuild incrementally: on `ART_PATCH` remove+insert that id; on `ART_CREATE/DELETE` insert/remove.
- Exclude non-visual artifacts if added later.

Operations:

- `queryPoint(p: Vec2): Id[]` → for click hit-test.
- `queryAABB(r: AABB): Id[]` → for marquee selection and snap candidate window.
- `nearestAnchors(p: Vec2, r: number): Anchor[]` → small radius search for snapping.

Libraries: `rbush` or `flatbush` acceptable; custom quadtree is trivial and predictable.

Anchor Index

Purpose: snapping and connector port resolution.

- For each artifact, precompute anchors: $\{type: 'center' | 'midL' | 'midR' | 'midT' | 'midB' | 'corner', pos: Vec2\}$.
- Store in a multimap `artifactId -> Anchor[]`.
- For connectors, expose route waypoints as anchors when needed.

Computation:

- Rebuild per artifact change; $O(1)$ per artifact (constant number of anchors).

Z-Order

- Maintain a flat order: `Id[]`.
- Hit-test resolves candidates from spatial index, then sort by order descending to get the topmost id.
- Bring-to-front/back operations mutate the array; keep $O(1)$ by tracking an index map if necessary.

Dirty Regions

- Each reducer returns a list of changed AABBs.
 - Renderer unions overlapping regions and redraws minimal tiles.
-

Interaction Pipeline

1. **Pointer input** (down/move/up) captured on Canvas.
2. **Hit-test**
 - `queryPoint` for primary; if multiple, choose topmost via z-order.
 - For connectors, include edge hit area by expanding segments to a thickness for picking.
3. **Gizmo**
 - Selection gizmo, transform handles, connector dragger.
 - Emits high-level intents: `dragTranslate(dx,dy)`, `resize(handle,dx,dy)`, `connect(fromId,port,toPoint)`.
4. **Snap**
 - If snap enabled, gather candidates with `queryAABB` around the moving AABB expanded by a radius.
 - Compute deltas to nearest anchors or grid lines; apply strongest rule (exact overlap > edge align > grid).
5. **Command**
 - Convert intent (+ snapped delta) into `Command`.
6. **Reducer**
 - Pure function updates `Sheet`.
7. **Derived**
 - Update quadtree/anchors/z-order indices incrementally.
8. **Render**
 - Redraw dirty regions.
9. **CRDT Sync**
 - Apply equivalent Yjs updates; broadcast.

Determinism preserved: same command sequence → same state.

Connectors and Routing (Stage)

- Orthogonal (Manhattan) only.
- Ports: infer side based on nearest side of source/target AABBs.
- Router algorithm: grid-A* over inflated obstacle set.
 - Inflate all artifact AABBs by connector margin.
 - Start from a source port point, A* through grid cells to target port area.
 - Post-process with segment simplification (remove collinear bends).
- Cache route per edge keyed by (fromGeom,toGeom,obstacleVersion).
- Recompute routes only when endpoints or nearby obstacles change (query obstacles via quadtree).

Dependencies: optional small pathfinding lib; otherwise custom A* is short.

Snapping

- Grid snap at `config.grid` interval.
 - Object snap to anchors from nearby artifacts (quadtree query radius = a few grid units).
 - Angle snap for connector creation (0/90/45°); implemented as clamped direction vectors.
 - “Sticky” constraints deferred; initial stage creates no persistent constraints.
-

Rendering Model (Canvas2D)

- Back→front draw using order.
 - Draw pipeline: grid → edges → artifacts → selection/handles → cursors.
 - Text measuring cache keyed by (font,size,text).
 - Mermaid/MDX rendered in offscreen DOM to SVG/canvas bitmap, then drawn into the board rectangle; cache rasterization per source hash.
 - Export to PNG/SVG: reuse same draw routines; SVG path generation mirrors Canvas calls.
-

Persistence Model (PostgreSQL)

Tables (minimal):

```
-- Project/sheet grouping
CREATE TABLE projects (
  id uuid PRIMARY KEY,
  name text NOT NULL,
  owner_id uuid NOT NULL,
  created_at timestamptz DEFAULT now()
```

```

);

CREATE TABLE sheets (
  id uuid PRIMARY KEY,
  project_id uuid REFERENCES projects(id) ON DELETE CASCADE,
  name text NOT NULL,
  config jsonb NOT NULL DEFAULT '{}',
  created_at timestamptz DEFAULT now()
);

-- Event log of CRDT deltas or normalized commands
CREATE TABLE sheet_events (
  id bigserial PRIMARY KEY,
  sheet_id uuid REFERENCES sheets(id) ON DELETE CASCADE,
  actor_id uuid NOT NULL,
  ts timestamptz NOT NULL DEFAULT now(),
  payload jsonb NOT NULL
);

-- Periodic materialized snapshots for fast load
CREATE TABLE sheet_snapshots (
  sheet_id uuid PRIMARY KEY REFERENCES sheets(id) ON DELETE CASCADE,
  version bigint NOT NULL,
  state jsonb NOT NULL,           -- normalized {artifacts, edges, order, config}
  updated_at timestamptz NOT NULL DEFAULT now()
);

```

Load strategy:

- On join: return latest snapshot + subsequent events cursor.
- Client reconstructs; CRDT merges incoming updates.

Save strategy:

- Append CRDT updates to sheet_events.
- Background job materializes sheet_snapshots every N events or T seconds.

API Surface (Stage)

HTTP (Next.js route handlers):

- GET /api/sheets/:id → { snapshot, eventsSince }.
- POST /api/sheets → create sheet.
- POST /api/sheets/:id/snapshot → force snapshot (admin/dev).
- GET /api/export/:sheetId.svg|png → stream export.

WebSocket (Node/Yjs server):

- `join { sheetId }`
- `update { yUpdate }`
- `presence { cursor, selection }`

Security:

- Cookie session for HTTP; WS token bearing sheet access.
-

Extensibility Hooks (Sheet Level)

- `registerArtifact(kind, { schema, measure, draw, anchors, hit, routerHooks? })`
 - `schema`: Zod validation for props.
 - `measure`: precompute sizes (e.g., MDX content).
 - `draw`: uses a target-agnostic API; adapter maps to Canvas/SVG.
 - `anchors`: returns anchor list for snapping/ports.
 - `hit`: specialized shape hit-test if not rectangular.
 - `routerHooks`: optionally expose obstacle geometry for connectors.
 - `registerTool(name, handlers)` for UI tools, independent from core.
 - `registerExport(name, fn(sheet)->Blob)`.
-

Compute Placement

Client:

- Reducers and indices.
- Quadtree, anchors, hit-testing.
- Routing for small/medium graphs.
- MDX/Mermaid compilation and rasterization.
- Export PNG/SVG (fast path).

Server:

- Optional heavy routing/layout batches.
 - Export PDF (headless renderer).
 - Periodic snapshot job.
 - AI agent endpoints that submit command batches.
-

Performance Targets

- 10k artifacts, 2k connectors at 60fps on modern laptop with Canvas2D.
 - Hit-test: $O(\log n)$ per pointer event via quadtree; worst-case small constant candidate set.
 - Drag frame latency $\leq 8\text{ms}$ at 120Hz for modest scenes; use dirty region redraw only.
-

Error/Recovery

- Deterministic reducers guarantee replay.
 - Snapshot mismatch \rightarrow client re-syncs from latest snapshot.
 - Corrupt event payloads rejected by Zod; server logs and drops.
-

Testing

- Reducers: property-based tests on command sequences.
 - Spatial index: randomized insert/remove/query invariants.
 - Router: golden tests on fixture boards.
 - Rendering: pixel-diff tests for primitives; skip platform font variance by using test font.
-

Milestone Breakdown (Sheet Model)

Week 1

- Core types and reducers.
- Quadtree + anchors + z-order.
- Canvas renderer (rect, text).
- Command bus with undo/redo.
- Grid + basic snapping.

Week 2

- Connectors + orthogonal router + edge hit-testing.
- Marquee select + grouping.
- MDX artifact (render \rightarrow bitmap).
- Mermaid artifact (render \rightarrow SVG \rightarrow bitmap).
- PNG/SVG export.

Week 3

- Yjs sync + presence.
- Postgres events + snapshots + loader.
- Snapshot cron.

- E2E tests on create/move/resize/connect/export.
 - Profiling + dirty-region optimization.
-

Justifications

- Client-first compute minimizes server cost and lowers perceived latency.
 - CRDT + event log yields robust collaboration and replayability.
 - Quadtree/anchor indices provide predictable, explainable performance.
 - Tagged-union artifacts keep type-level safety with straightforward extensibility.
 - JSONB storage matches evolving schemas without migrations on every artifact tweak.
 - Canvas2D is sufficient for stage performance; WebGL is a drop-in renderer later.
 - MDX/Mermaid as first-class artifacts validate the “content + flow” model early without server complexity.
-

Open Space (“Thinking Room”)

- Persistent constraints (Cassowary/Kiwi) for true layout rules beyond snap.
- Id-buffer color picking if Canvas hit precision becomes a bottleneck.
- Incremental route graph (channel routing) to avoid per-edge A*.
- Multi-sheet projections sharing a single underlying graph.
- Pluggable dataflow execution graph for computed artifacts.
- On-sheet code artifacts with sandboxed evaluators.
- Semantic merges on artifacts beyond CRDT structural merges.
- WebGPU renderer for very large scenes if needed later.