

Here's the breakdown of an initial architecture that fits your constraints (cheap, quick, declarative, algorithm-driven, extendable), and explicitly answers where each responsibility sits.

1. Platform Decision

Yes, Next.js is a good path — but run it in a *client-heavy, server-light* mode.

- Board, flow, and artifact logic stay in the browser to reduce server compute.
- Server focuses on persistence, collaboration, and controlled compute tasks.

For development:

- One repo, pnpm workspaces.
 - Local Postgres via Docker.
 - Node server co-located with Next.js during dev.
 - Deployment later can split them into pods in k8s.
-

2. Roles of Each Layer

Next.js Client

- Canvas rendering of sheets and artifacts.
- Interaction, selection, snapping, layout.
- Declarative model management in memory (CRDT-backed).
- Local execution of lightweight algorithms (layout, routing, spatial index).
- Temporary artifact state before persisting.
- MDX rendering and preview.
- Mermaid diagram rendering (client-side library).

Next.js Server (API routes)

- CRUD for sheets, artifacts, users.
- Auth and session control.
- Validation of incoming mutations.
- Export endpoints (SVG/PNG/PDF/MDX).
- Metadata search and filters.
- Hooks for AI agents to inject workflows.

Dedicated Node Server (or same process in dev)

- Real-time sync server (Yjs WebSocket or custom).
- Long-running compute:

- Graph layout for large flows.
 - Type/schema generation from artifacts.
 - AI agent orchestration (workflow suggestions, auto-placement).
- Webhooks/event bus for external integrations.
-

3. Data Model (Essentials)

Tables in Postgres

- projects: id, name, owner_id, created_at.
- sheets: id, project_id, name, config, created_at.
- artifacts: id, sheet_id, type, props (JSONB), geom (JSONB), style (JSONB), meta (JSONB).
- edges: id, sheet_id, from_artifact, to_artifact, type, props (JSONB).
- events: id, sheet_id, actor_id, type, payload (JSONB), ts.
- users: id, email, hashed_password.

Artifact Type Examples

- "box": geom, fill, label.
 - "mdx": source, compiled_ast.
 - "mermaid": source, compiled_svg.
 - "tasklist": tasks array, completion state.
-

4. State and Flow

In-Memory State

- Client keeps a CRDT document (per sheet).
- CRDT merges remote + local changes.
- Reducers update immutable state; render system listens for diffs.

Persistence Flow

1. Client commits change to CRDT.
 2. CRDT sync server broadcasts change to peers.
 3. Server batches CRDT updates → Postgres as event log.
 4. Snapshots generated periodically for faster load.
-

5. Routes and Controllers

REST/GraphQL (pick one; REST easier to start)

- GET /api/projects/:id → project meta + sheet list.
- GET /api/sheets/:id → latest snapshot + events.
- POST /api/sheets → create sheet.
- PATCH /api/sheets/:id → update meta/config.
- POST /api/artifacts → add artifact.
- PATCH /api/artifacts/:id → update artifact props/geom/style.
- DELETE /api/artifacts/:id.
- GET /api/export/:sheet_id.svg|png|pdf|mdx.

Realtime (WebSocket)

- joinSheet(sheetId) → receive state + presence.
 - update(delta) → broadcast CRDT changes.
 - presence(cursor, selection) → live cursors.
-

6. Rendering on Sheet

Client-Side Algorithms

- Quadtree for hit testing.
- Constraint solver (Kiwi.js) for snapping/alignment.
- A* for connector routing.
- Layout algorithms (DAG, force-directed) from lightweight libs (elkjs, dagre-lite).
- Markdown → MDX → React tree via unified pipeline.
- Mermaid parser → SVG renderer in browser.

Artifacts render in a z-indexed scene:

- Base primitives (rect, ellipse, line).
 - Embeds (mdx, mermaid, image).
 - Interaction layer (selection boxes, transform handles).
-

7. Storage Strategy

- **Postgres**: primary store (JSONB for flexible props).
 - **S3-compatible bucket**: images, large attachments.
 - **Local IndexedDB**: offline board state + assets cache.
-

8. Algorithms from pnpm Ecosystem

- **CRDT**: yjs + y-websocket.
 - **Layout**: elkjs (WASM) or dagre-lite for graphs.
 - **Routing**: custom A* or pathfinding npm for grids.
 - **Spatial index**: rbush or flatbush.
 - **Constraint solver**: kiwi (Cassowary).
 - **MDX**: @mdx-js/mdx, remark-*, rehype-*.
 - **Mermaid**: mermaid npm (tree-shakable).
 - **Schema validation**: zod.
 - **DB**: drizzle-orm with Postgres driver.
-

9. Execution Split

Computed in Client

- Interactive transforms.
- Snapping, hit-testing.
- Rendering, MDX compilation, Mermaid rendering.
- Small/medium graph layouts.

Computed in Node Server

- Large graph layouts (batch).
 - Export rendering (high-res SVG/PDF).
 - AI agent workflows.
 - Batch type/schema generation.
-

10. Initial Roadmap

1. **Core sheet model**: CRDT doc, artifact base type, quadtree.
 2. **Rendering engine**: Canvas2D with base shapes.
 3. **Persistence**: Postgres schema + simple REST.
 4. **Realtime**: y-websocket integration.
 5. **Artifact plugins**: box, mdx, mermaid, tasklist.
 6. **Export pipeline**: SVG/PNG.
 7. **Extensibility hooks**: registerArtifact, registerLayout.
 8. **AI hooks**: endpoint to inject CRDT updates.
-

This approach lets you build from scratch with clear separation of compute and storage, keeps the initial implementation cheap and fast, and positions the sheet/artifact model to extend into complex workflows without refactoring the core.